

BorgBackup

Mohammad Fazeli

May 19, 2021

"I found the Holy Grail of backups."
(Stavros K. about Attic-Backup, 8/2013)

Backup Desiderata

- ▶ Take very little space
- ▶ Encrypted
 - ▶ Backup place hacked means everything is gone
- ▶ Multiple versions of data
- ▶ Pruning versions By Grandfather-Father-Son
- ▶ Deduplicated, reduced space

Backup Desiderata: Not using Borg

- ▶ Disk failure shouldn't cause data loss
- ▶ Server failure shouldn't effect availability
- ▶ Disks have bit-rot

My Solution

Borg + BTRFS RAID1 on at least two locations BTRFS RAID1
removes bit-rot problem two locations solves availability

What is BorgBackup

BorgBackup (short: Borg) deduplicating backup program (compressed and encrypted authenticated)

Borg Features

- ▶ cli backup tool, separate GUIs
- ▶ good arch / platform / fs support
- ▶ deduplication
- ▶ compression
- ▶ auth. encryption
- ▶ read from locally mounted fs
- ▶ store to local fs or to remote borg
- ▶ Fuse-mount backup archives
- ▶ HW-accelerated crypto

Borg as a Project

- ▶ community project
- ▶ FOSS (BSD license)
- ▶ Python 3.7(current, 3.6-9 supported)
- ▶ for speed: a little Cython & C
- ▶ good docs (for users, devs)
- ▶ good test coverage(83% today), travis CI
- ▶ github: [borgbackup/community](https://github.com/borgbackup/community)
- ▶ Pays money for security bugs(10,500\$ paid today)

In Practice

- ▶ size of postgres db
- ▶ 48 hourly
- ▶ 20 daily
- ▶ Weekly 6
- ▶ Monthly ...
- ▶ 78 instances each 2GB
- ▶ all:153.06 GB
- ▶ compressed 41.24 GB
- ▶ deduplicated 22.50 GB

In Practice

- ▶ media folder of a big django project
- ▶ 3 instances each about 63 GB
- ▶ 182.70 GB compressed all instances (incompressible mostly already compressed images)
- ▶ 60.13 GB deduplicated across time

In Practice

- a little changing postgres db
- 12 instances
- 7 daily
- 4 weekly
- 1 monthly
- 2.75 GB all
- compressed 180MB
- deduplicated 30 MB

In Practice

- a wordpress site files
- 32 instances
- 80 GB total
- 70 GB compressed (incompressible mostly images)
- deduplicated 2.68 GB

In Practice

- a wordpress site db
- 32 instances
- 5.5 GB
- 420 MB compressed
- 204 MB deduped

In Practice

- a lot of mail boxes
- 20 instances
- 20 daily
- total 727GB
- 348 GB compressed (compressible since it is text)
- deduped 17.5 GB

In Practice

Total size: 1145 GB

Fitted in 103GB

It scales well

one of developers of borg
ssh://borg@myserver/repos/myrepo
Original size 22.76 TB
Compressed size 18.22 TB
Dedup size 486.20 GB
Unique chunks 6305006
Total chunks 272643223

Where to put

Own server with ssh and free space:
install borg, configure, done!

No own remote server? No problem:

- ▶ rsync.net (ssh, cli)
- ▶ hetzner's storage box
- ▶ borgbase.com (ssh, web, easy)
- ▶ local repo + rclone to cloud

Deduplicating

dedup does **NOT** depend on:

- ▶ file/directory names:
 - ▶ move file even between machines, no change required
- ▶ Whole file hash and timestamps:
 - ▶ only a small part of a file changed, only the change, VMs , or raw disks, only changes are synced.
- ▶ Place of chunk in a file:
 - ▶ stuff moved around inside a file , doesn't increase it.

Encryption

Enc :

- ▶ Tampering/Corruption detection by HMAC or blake2b
- ▶ 256 bit AES- CTR
- ▶ uses OpenSSL(libcrypto)
- ▶ Passphrase and a key, could be separate
- ▶ PDKDF2 100K rounds

Compression

- ▶ none
- ▶ lz4: fast, low compression(faster than none :-|)
- ▶ zstd: from high speed low compression to slow and high compression
- ▶ zlib: medium speed and compression
- ▶ lzma: low speed, high compression

Safe

borg uses:

- checksums
- transactions
- fs: syncing, atomic ops
- backup repo- log-like KV store
- checkpoint while backing up

language

```
# initialize a repository:  
borg init /tmp/borg
```

```
# create a "first" archive inside this repo (verbose)  
borg create --progress --stats\  
    /tmp/borg::first ~/Desktop
```

```
# create a "second" archive (less verbose):  
borg create /tmp/borg::second ~/Desktop
```

```
# even more verbose:  
borg create -v --stats /tmp/borg::third ~/Desktop
```

language

```
# list repo / archive contents:
```

```
borg list /tmp/borg
```

```
borg list /tmp/borg::first
```

```
# extract ("restore") from an archive to cwd:
```

```
mkdir test ; cd test
```

```
borg extract /tmp/borg::third
```

```
# simulate extraction (good test):
```

```
borg extract -v --dry-run /tmp/borg::third
```

```
# check consistency of repo:
```

```
borg check /tmp/borg
```

language

```
^^|^^|# list repo / archive contents:
^^|^^|
^^|^^|borg list /tmp/borg
^^|^^|borg list /tmp/borg::first
^^|^^|
^^|^^|# extract ("restore") from an archive to cwd:
^^|^^|
^^|^^|mkdir test ; cd test
^^|^^|borg extract /tmp/borg::third
^^|^^|
^^|^^|# simulate extraction (good test):
^^|^^|borg extract -v --dry-run /tmp/borg::third
^^|^^|
^^|^^|# check consistency of repo:
^^|^^|borg check /tmp/borg
^^|
```